

```
1 #include <stdio.h>
2
3 // cuda headers
4 #include <cuda.h>
5 #include "helper_timer.h"
6
7 // macros
8 #define BLOCK_WIDTH 32
9
10 // global variables
11 int *hostA=NULL;
12 int *hostB=NULL;
13 int *hostC=NULL;
14 int *gold=NULL;
15
16 int *deviceA=NULL;
17 int *deviceB=NULL;
18 int *deviceC=NULL;
19
20 float timeOnCPU = 0.0f;
21 float timeOnGPU = 0.0f;
22
23 // cuda kernel function
24 __global__ void matMulGPU(int *A,int *B,int *C,int numRows,int
numAColumns,int numBColumns,int numCColumns)
25 {
26     // variable declarations
27     int row=blockIdx.y * blockDim.y + threadIdx.y;
28     int column=blockIdx.x * blockDim.x + threadIdx.x;
29     // code
30     if((row < numRows) && (column < numBColumns))
31     {
32         int value=0.0;
33         for(int k=0; k < numAColumns; k++)
34         {
35             int a=A[row * numAColumns + k];
36             int b=B[k * numBColumns + column];
37             value += a*b;
38         }
39         C[row * numCColumns + column]=value;
40     }
41 }
42
43 int main(int argc,char *argv[])
44 {
45     // function declarations
46     void InitA(int *data,int,int);
47     void InitB(int *data,int,int);
48     void matMulCPU(int*, int*, int*, int, int, int, int);
49     void cleanup(void);
50
51     // variable declarations
52     int numRows=BLOCK_WIDTH;
53     int numAColumns=BLOCK_WIDTH;
54     int numBRows=BLOCK_WIDTH;
55     int numBColumns=BLOCK_WIDTH;
```

```
56
57     int numCRows=numARows;
58     int numCColumns=numBColumns;
59
60     int numGoldRows=numARows;
61     int numGoldColumns=numBColumns;
62
63     int sizeA = numARows * numAColumns * sizeof(int);
64     int sizeB = numBRows * numBColumns * sizeof(int);
65     int sizeC = numCRows * numCColumns * sizeof(int);
66     int sizeGold = numGoldRows * numGoldColumns * sizeof(int);
67
68     cudaError_t result = cudaSuccess;
69
70     // code
71     // host memory allocation
72     hostA=(int *)malloc(sizeA);
73     if(hostA==NULL)
74     {
75         printf("Host Memory allocation is failed for hostA matrix.\n");
76         cleanup();
77         exit(EXIT_FAILURE);
78     }
79
80     hostB=(int *)malloc(sizeB);
81     if(hostB==NULL)
82     {
83         printf("Host Memory allocation is failed for hostB matrix.\n");
84         cleanup();
85         exit(EXIT_FAILURE);
86     }
87
88     hostC=(int *)malloc(sizeC);
89     if(hostC== NULL)
90     {
91         printf("Host Memory allocation is failed for hostC matrix.\n");
92         cleanup();
93         exit(EXIT_FAILURE);
94     }
95
96     gold=(int *)malloc(sizeGold);
97     if(gold== NULL)
98     {
99         printf("Host Memory allocation is failed for gold matrix.\n");
100         cleanup();
101         exit(EXIT_FAILURE);
102     }
103
104     // printing matrix dimensions and sizes
105     printf("The Dimensions Of Matrix 'hostA' Are : %d x %d\n",numARows,numAColumns);
106     printf("The Dimensions Of Matrix 'hostB' Are : %d x %d\n",numBRows,numBColumns);
107     printf("The Dimensions Of Matrix 'hostC' Are : %d x %d\n",numCRows,numCColumns);
108
```

```
109 printf("The Dimensions Of Matrix 'gold' Are : %d x %d\n",numGoldRows,numGoldColumns);
110
111 printf("Size Of Matrix hostA = %d\n",sizeA);
112 printf("Size Of Matrix hostB = %d\n",sizeB);
113 printf("Size Of Matrix hostC = %d\n",sizeC);
114
115 printf("Size Of Matrix gold = %d\n",sizeGold);
116
117 // fill source matrices
118 InitA(hostA,numARows,numAColumns);
119 InitB(hostB,numBRows,numBColumns);
120
121 // device memory allocation
122 result=cudaMalloc((void **)&deviceA,sizeA);
123 if(result!=cudaSuccess)
124 {
125     printf("Device Memory allocation is failed for deviceA matrix.\n");
126     cleanup();
127     exit(EXIT_FAILURE);
128 }
129
130 result=cudaMalloc((void **)&deviceB,sizeB);
131 if(result!=cudaSuccess)
132 {
133     printf("Device Memory allocation is failed for deviceB matrix.\n");
134     cleanup();
135     exit(EXIT_FAILURE);
136 }
137
138 result=cudaMalloc((void **)&deviceC,sizeC);
139 if(result!=cudaSuccess)
140 {
141     printf("Device Memory allocation is failed for deviceC matrix.\n");
142     cleanup();
143     exit(EXIT_FAILURE);
144 }
145
146 // copy data from host matrices into device matrices
147 result=cudaMemcpy(deviceA,hostA,sizeA,cudaMemcpyHostToDevice);
148 if(result!=cudaSuccess)
149 {
150     printf("Host to Device Data Copy is failed for deviceA matrix.\n");
151     cleanup();
152     exit(EXIT_FAILURE);
153 }
154
155 result=cudaMemcpy(deviceB,hostB,sizeB,cudaMemcpyHostToDevice);
156 if(result!=cudaSuccess)
157 {
158     printf("Host to Device Data Copy is failed for deviceB matrix.\n");
159     cleanup();
160     exit(EXIT_FAILURE);
161 }
162
163 // CUDA kernel configuration
```

```
164 dim3 dimGrid=dim3(ceil((int)numBColumns/(int)BLOCK_WIDTH),ceil((int)  ?
    numRows/(int)BLOCK_WIDTH),1);
165 dim3 dimBlock=dim3(BLOCK_WIDTH,BLOCK_WIDTH,1);
166
167 // CUDA kernel for matrix multiplication
168 StopwatchInterface* timer = NULL;
169 sdkCreateTimer(&timer);
170 sdkStartTimer(&timer);
171
172 matMulGPU <<<dimGrid, dimBlock >>>(deviceA, deviceB, deviceC, numRows,  ?
    numAColumns, numBColumns, numCColumns);
173
174 sdkStopTimer(&timer);
175 timeOnGPU = sdkGetTimerValue(&timer);
176 sdkDeleteTimer(&timer);
177 timer = NULL;
178
179 // copy data from device matrix into host matrix
180 result=cudaMemcpy(hostC,deviceC,sizeC,cudaMemcpyDeviceToHost);
181 if(result!=cudaSuccess)
182 {
183     printf("Device to Host Data Copy is failed for hostC matrix.\n");
184     cleanup();
185     exit(EXIT_FAILURE);
186 }
187
188 // matrix multiplication on host
189 matMulCPU(hostA, hostB, gold, numRows, numAColumns, numBColumns,  ?
    numCColumns);
190
191 // comparison
192 int breakValue = -1;
193 bool bAccuracy = true;
194 for (int i = 0; i < numRows * numCColumns; i++)
195 {
196     int val1 = gold[i];
197     int val2 = hostC[i];
198     if (val1 != val2)
199     {
200         bAccuracy = false;
201         breakValue = i;
202         break;
203     }
204 }
205
206 char str[128];
207 if (bAccuracy == false)
208     sprintf(str, "Comparison of CPU and GPU Matrix Multiplication is not  ?
        accurate at array index %d", breakValue);
209 else
210     sprintf(str, "Comparison of CPU and GPU Matrix Multiplication is  ?
        accurate");
211
212 printf("Time taken for Matrix Multiplication on CPU = %.6f\n", timeOnCPU);
213 printf("Time taken for Matrix Multiplication on GPU = %.6f\n", timeOnGPU);
214 printf("%s\n", str);
```

```
215
216     // cleanup
217     cleanup();
218
219     return(0);
220 }
221
222 void InitA(int *data,int row,int col)
223 {
224     int num=1;
225     // code
226     for(int i=0;i<row;i++)
227     {
228         for(int j=0;j<col;j++)
229         {
230             *(data + i * col + j)=num;
231             num++;
232         }
233     }
234 }
235
236 void InitB(int *data,int row,int col)
237 {
238     int num=BLOCK_WIDTH;
239     // code
240     for(int i=0;i<row;i++)
241     {
242         for(int j=0;j<col;j++)
243         {
244             *(data + i * col + j)=num;
245             num--;
246         }
247     }
248 }
249
250 void matMulCPU(int* A, int* B, int* C, int numRows, int numAColumns, int numBColumns, int numCColumns)
251 {
252     // code
253     StopwatchInterface* timer = NULL;
254     sdkCreateTimer(&timer);
255     sdkStartTimer(&timer);
256
257     for (int i = 0; i < numRows; ++i)
258     {
259         for (int j = 0; j < numBColumns; ++j)
260         {
261             int value = 0.0f;
262             for (int k = 0; k < numAColumns; ++k)
263             {
264                 int a = A[i * numAColumns + k];
265                 int b = B[k * numBColumns + j];
266                 value += a * b;
267             }
268             C[i * numCColumns + j] = value;
269         }
270     }
```

```
270     }
271
272     sdkStopTimer(&timer);
273     timeOnCPU = sdkGetTimerValue(&timer);
274     sdkDeleteTimer(&timer);
275     timer = NULL;
276 }
277
278 void cleanup(void)
279 {
280     // code
281     if (deviceC)
282     {
283         cudaFree(deviceC);
284         deviceC = NULL;
285     }
286
287     if (deviceB)
288     {
289         cudaFree(deviceB);
290         deviceB = NULL;
291     }
292
293     if (deviceA)
294     {
295         cudaFree(deviceA);
296         deviceA = NULL;
297     }
298
299     if (gold)
300     {
301         free(gold);
302         gold = NULL;
303     }
304
305     if (hostC)
306     {
307         free(hostC);
308         hostC = NULL;
309     }
310
311     if (hostB)
312     {
313         free(hostB);
314         hostB = NULL;
315     }
316
317     if (hostA)
318     {
319         free(hostA);
320         hostA = NULL;
321     }
322 }
323
```